

PARAG K. GOYAL

+91 9602563128 | 1997parag.goyal@gmail.com | linkedin.com/in/parag-k-goyal

TECHNICAL SKILLS

Languages: Java, JavaScript

Backend: Spring Boot, Spring MVC, Micronaut

Frontend: React

Data & Search: Oracle Autonomous Transaction Processing (ATP), MySQL, OpenSearch

Tools: Git, Maven, Swagger/OpenAPI, Postman, Mockito, JUnit, Linux

PROFESSIONAL EXPERIENCE

Senior Application Software Engineer | Oracle

Sep 2024 – Present

- Delivered a scalable Micronaut-based REST API layer for Oracle Health's clinical charting service, enabling system-configured forms to be assembled, served, and persisted for practitioners, by designing and building out DAO, service, DTO, and controller layers integrated with Oracle Vector Database.
- Reduced duplicate data entry for practitioners, giving them policy-controlled access to prior values based on global, practitioner, time-window, and element-level settings, by enabling cross-form retrieval and auto-population of last-charted values for shared elements.
- Restricted charting-form actions to authorized clinicians only, securing sign, get, save, and update operations, by designing and implementing policy-level authorization checks across those endpoints.
- Enabled downstream reuse of structured clinical data across text, rich text, date/time, numerical, and single-/multi-select formats to support clinician sign workflows, by leading the REST integration with Oracle Health's Observation service.
- Ensured the UI always rendered an accurate, point-in-time view of a form even though its sections and questions were independent, shared, editable objects, by building an API that recursively resolved a form into a snapshot; optimized generation with a Java multithreading approach and a hashmap lookup keyed on a CRC-based signature hash of the form's sections and questions, so unchanged forms reused an existing snapshot instead of recomputing one.
- Improved paginated search performance by 20%, by migrating the query path from Oracle Vector Database to Oracle ATP and implementing token-based pagination, targeted indexing, and parameterized queries.
- Improved metadata discoverability for an AI-assisted form builder, by building an OpenSearch ingestion pipeline that normalized questionnaire, section, and element-definition metadata.
- Strengthened release confidence across two production releases, serving as the release automation point of contact and reporting execution metrics to the test automation team, by developing and integrating end-to-end automated tests covering key API workflows into the clinical charting service pipeline.
- Shortened the ramp-up time for new engineers to reach productive contribution, by onboarding them to the clinical charting service's layered architecture and save/sign workflow implementation.
- Maintained service reliability in production, resolving issues as they arose, by providing on-call diagnostic and troubleshooting support.

Software Engineer | Brane Enterprises

Mar 2023 – Sep 2024

- Gave clients a way to configure data workflows — including SELECT, JOIN, and aggregation operations — without persisting their data, by developing an adapter framework for a low-code/no-code integration platform that processed client data at runtime.
- Expanded the range of client-configurable workflows to include stored-procedure execution, by extending the JDBC adapter framework.

Software Development Engineer I | Amazon

Aug 2021 – Mar 2023

- Guided customers toward similar product offerings across multiple marketplaces with localized messaging, coordinating with product managers and using Amazon's internal translation service with marketplace context, by developing and deploying the Featured Offer Disqualified (FOD) widget for Amazon product detail pages.
- Improved alternative-offer coverage on product listings when no Featured Offer was available, by authoring design options for a Product Recommendation widget in collaboration with the recommendation team.

Specialist Programmer | Infosys

Aug 2019 – Aug 2021

- Automated processing of client-defined input messages into persisted, enriched database records, by building a RabbitMQ-based processor that extracted relevant data and derived additional properties.
- Eliminated manual data entry for product creation from structured XML data, including text and images, by building an XML parser and a cron-scheduled background job.

PROJECTS

WhizChat | Java, Spring Boot, PostgreSQL, JPA, WebSocket/STOMP | GitHub

- Built a real-time messaging application with Spring Boot REST APIs, JPA/PostgreSQL persistence, Keycloak JWT authentication, and WebSocket/STOMP notifications; implemented chats, media uploads, read receipts, and per-user notifications.

Movies | React, JavaScript, Spring Boot, MongoDB | GitHub

- Built a full-stack movie management application using React, JavaScript, Spring Boot, and MongoDB, with responsive views, REST API integration, trailer playback, and review submission.

EDUCATION

Malaviya National Institute of Technology (MNIT), Jaipur

B.Tech., 2015 – 2019